

Aura by Centro

Enterprise AI Co-Pilot

Helping businesses drive meaningful change for growth.


Local Gemma LLM

Multi-tenant & sandboxed

RAG + CAG memory


Zoho · Odoo · Genesys (MCP)

CLIENTS



A Advanced & Agentic RAG — beyond a naive chatbot

Messy human question → optimized query → sandboxed, reranked context → grounded answer with ■ citations



IMPACT Exact-match + semantic recall · sharply reduced hallucinations · RBAC enforced on every path · graceful degradation if any component is offline.
All pillars are config-flagged in backend/config.py and implemented in core/{chunking,retrieval,rerank,agent}.py.

A Why Aura wins the room

Not a wrapper

Hybrid + rerank + agentic retrieval — a defensible architecture, not a thin API shim.

Private & compliant

Every byte stays on Centro infra; tenant isolation enforced at the DB index.

Safe by design

Grounded-only answers, no silent writes, full audit trail + CSV export.

Cheap & scalable

\$0 software licensing; one config flip from laptop to a multi-GPU AWS fleet.

A Where we are & what's next

DELIVERED (working POC)

- ✓ Multi-tenant vector sandboxing (zero data leakage)
- ✓ Semantic CAG cache — instant, LLM-bypassing answers
- ✓ Dual-confirmation Action Cards for all writes
- ✓ Dynamic API schemas (no prompt hardcoding)
- ✓ Admin dashboard + self-service document uploads
- ✓ Embeddable in Zoho People; branded desktop client
- ✓ Runs on local Gemma — \$0 software licensing

NEXT (to production)

- > **Bigger + multimodal model**
Gemma 3 12B/27B or vision model on AWS GPU
- > **Production SSO**
Centro identity issues scoped tokens
- > **Live MCP integrations**
Real Odoo/Zoho/Genesys writes
- > **Desktop packaging**
Electron / Tauri / PWA — decision pending
- > **Scale on AWS**
vLLM GPU fleet sized for 1,500 users